

Training Deep Nets Using Conjugate Deep Nets

The rate of progress of gradient descent on a function $F(x)$ depends on the shape of F : The closer ∇F points to the minimizer of F , the faster the progress. Preconditioning is the process of aligning ∇F with the minimizer by introducing a change of variables $x = G(y)$, and performing gradient descent on $F(G(y))$ instead. The solution y^* is then mapped back via $x^* = G^{-1}(y^*)$. Typically, G is chosen to cause the Hessian $\nabla^2 F(G(y))$ to be close to identity. But any change of variables that causes $\nabla F(G(y))$ to point along $y - y^*$ would work. This document proposes one such preconditioner, based on the stochastic Hessian.

Notation

We'll write a deep net as a pipeline of functions $\ell = 1, \dots, L$,

$$z_1 = f_1(z_0; x_1) \tag{1}$$

$$\dots \tag{2}$$

$$z_\ell = f_\ell(z_{\ell-1}; x_\ell) \tag{3}$$

$$\dots \tag{4}$$

$$z_L = f_L(z_{L-1}; x_L) \tag{5}$$

The vectors x_1, \dots, x_L are the parameters of the pipeline. The vectors z_1, \dots, z_L are its intermediate activations, and z_0 is the input to the pipeline. The last layer f_L computes the final activations and their training loss, so the scalar z_L is the loss of the model on the input z_0 . To make this loss's dependence on z_0 and the parameters explicit, we'll sometimes write it as $z_L(z_0; x)$. The foregoing formalization deviates slightly from the traditional deep net formalism in two ways: First, the training labels are subsumed in z_0 , and are propagated through the layers until they're used in the loss. Second, the last layer fuses the loss (which has no parameters) and the last layer (which does).

Backpropagation and forward propagation, the recursive way

We would like to fit the vector of parameters $x = (x_1, \dots, x_L)$ given a training dataset, which we'll represent by a stochastic input z_0 to the pipeline. Training the model proceeds by gradient descent steps along the stochastic gradient $\partial z_L(z_0; x) / \partial x$. The components of this direction can be computed by the chain rule with a backward recursion:

$$\begin{aligned} \frac{\partial z_L}{\partial x_\ell} &= \underbrace{\frac{\partial z_L}{\partial z_\ell} \frac{\partial z_\ell}{\partial x_\ell}}_{b_\ell^\top \nabla_x f_\ell^\top} \\ \frac{\partial z_L}{\partial z_\ell} &= \underbrace{\frac{\partial z_L}{\partial z_{\ell+1}} \frac{\partial z_{\ell+1}}{\partial z_\ell}}_{b_{\ell+1}^\top \nabla_z f_{\ell+1}^\top} \end{aligned}$$

The identification $b_\ell \equiv \frac{\partial z_L}{\partial z_\ell}^\top$, $\nabla_x f_\ell \equiv \frac{\partial z_L}{\partial x_\ell}^\top$, and $\nabla_z f_\ell \equiv \frac{\partial z_L}{\partial z_{\ell+1}}^\top$ turns this recurrence into

$$\begin{aligned}\frac{\partial z_L}{\partial x_\ell}^\top &= \nabla_x f_\ell \cdot b_\ell \\ d_\ell &= \nabla_z f_{\ell+1} \cdot b_{\ell+1}\end{aligned}$$

with the base case $b_L = 1$, a scalar.

Backpropagation and forward propagation, the matrix way

The above equations can be written in vector form as

$$\begin{aligned}\begin{bmatrix} \frac{\partial z_L}{\partial x_1} & \dots & \frac{\partial z_L}{\partial x_L} \end{bmatrix}^\top &= \underbrace{\begin{bmatrix} \nabla_x f_1 & & \\ & \ddots & \\ & & \nabla_x f_L \end{bmatrix}}_{\equiv D} \begin{bmatrix} b_1 \\ \vdots \\ b_L \end{bmatrix} \\ \underbrace{\begin{bmatrix} I & -\nabla_z f_2 & & \\ & I & -\nabla_z f_3 & \\ & & \ddots & \\ & & & I & -\nabla_z f_L \\ & & & & 1 \end{bmatrix}}_{\equiv M} \begin{bmatrix} b_1 \\ b_2 \\ b_3 \\ \vdots \\ b_{L-1} \\ b_L \end{bmatrix} &= \underbrace{\begin{bmatrix} 0 \\ \vdots \\ 1 \end{bmatrix}}_{\equiv e_L},\end{aligned}$$

or equivalently,

$$\begin{bmatrix} \frac{\partial z_L}{\partial x_1} & \dots & \frac{\partial z_L}{\partial x_L} \end{bmatrix}^\top = DM^{-1}e_L.$$

The matrix M is block bi-diagonal. Its diagonal entries are identity matrices, and its off-diagonal matrices are the gradient of the intermediate activations with respect to the layer's parameters. The matrix D is block diagonal, with the block as the derivative of each layer's activations with respect to its inputs. M is invertible because the spectrum of a triangular matrix can be read off its diagonal, which in this case is all ones.

The Hessian

The gradient we computed above is the unique vector v such that $dz_L \equiv z_L(x + dx) - z_L(x) \rightarrow v(x)^\top dx$ as $dx \rightarrow 0$. In this section, we'll compute the Hessian H of z_L with respect to the parameters. This is the unique matrix H that satisfies $dv \equiv v(x + dx) - v(x) \rightarrow H dx$ as $dx \rightarrow 0$. We'll use the facts that $dM^{-1} = -M^{-1}(dM)M^{-1}$ and $b = M^{-1}e_L$ to write

$$dv = d(DM^{-1}e_L) \tag{6}$$

$$= (dD)M^{-1}e_L + D(dM^{-1})e_L \tag{7}$$

$$= (dD)b - DM^{-1}(dM)M^{-1}e_L \tag{8}$$

$$= (dD)b - DM^{-1}(dM)b \tag{9}$$

We'll compute each of these terms separately.

As part of this agenda, we'll need to define the gradient of tensor-valued functions $g : \mathbb{R}^d \rightarrow \mathbb{R}^{o_1 \times \dots \times o_k}$. We'll define this gradient $\nabla_x g(x) \in \mathbb{R}^{(o_1 \dots o_k) \times d}$ as the unique matrix-valued function that satisfies $\text{vec}(g(x + dx) - g(x)) \rightarrow \nabla_x g(x) dx$ as $dx \rightarrow 0$. This convention allows us to readily define, for example, the Hessian of a vector-valued function: If $g : \mathbb{R}^d \rightarrow \mathbb{R}^o$, then $\nabla_{xx} g(x) \in \mathbb{R}^{o \times d^2}$ is the unique matrix such that $\text{vec}(\nabla_x g(x + dx) - \nabla_x g(x)) \rightarrow \nabla_{xx} g(x) dx$. This convention also supports the chain rule as expected: For example, the gradient of $h(x) \equiv f(g(x))$ for a matrix-valued f and g can be written as $\nabla f \nabla g$ as expected. The chain rule in turn lets us define mixed partial derivatives, like $\nabla_{yz} g$ for $g : \mathbb{R}^{|x|} \rightarrow \mathbb{R}^{|g|}$. For example, if $y \in \mathbb{R}^{|y|}$ and $z \in \mathbb{R}^{|z|}$ are restrictions of $x \in \mathbb{R}^{|x|}$ to some $|y|$ and $|z|$ -dimensional subsets, then $\nabla_z g(x) \in \mathbb{R}^{|g| \times |z|}$, and $\nabla_{yz} g(x) = \nabla_y \nabla_z g(x) \in \mathbb{R}^{|g| \times |z| \times |y|}$. See Chapter 6 of Magnus & Neudecker for a deeper treatment of higher order derivatives of vector-valued functions.

The term involving dD

The matrix D is diagonal with the ℓ th block containing the matrix $D_\ell \equiv \nabla_x f_\ell$. Using the fact that $\text{vec}(ABC) = (C^\top \otimes A) \text{vec}(B)$, we get

$$(dD)b = \begin{bmatrix} dD_1 & & \\ & \ddots & \\ & & dD_L \end{bmatrix} \begin{bmatrix} b_1 \\ \vdots \\ b_L \end{bmatrix} = \begin{bmatrix} dD_1 b_1 \\ \vdots \\ dD_L b_L \end{bmatrix} \quad (10)$$

$$= \begin{bmatrix} (b_1^\top \otimes I) \text{vec}(dD_1) \\ \vdots \\ (b_L^\top \otimes I) \text{vec}(dD_L) \end{bmatrix} \quad (11)$$

$$= \begin{bmatrix} b_1^\top \otimes I & & \\ & \ddots & \\ & & b_L^\top \otimes I \end{bmatrix} \begin{bmatrix} \text{vec}(dD_1) \\ \vdots \\ \text{vec}(dD_L) \end{bmatrix} \quad (12)$$

Observe that $\text{vec} dD_\ell = d\text{vec} \nabla_x f_\ell(z_{\ell-1}; x_\ell)$, varies with dx through both its arguments x_ℓ and $z_{\ell-1}$. Using mixed partials of vector-valued functions described above, we get

$$\text{vec}(dD_\ell) = d\text{vec}(\nabla_x f_\ell) = (\nabla_{xx} f_\ell) dx_\ell + (\nabla_{zx} f_\ell) dz_{\ell-1}.$$

Stacking these equations gives

$$\begin{bmatrix} \text{vec}(dD_1) \\ \vdots \\ \text{vec}(dD_L) \end{bmatrix} = \begin{bmatrix} \nabla_{xx} f_1 & & \\ & \ddots & \\ & & \nabla_{xx} f_L \end{bmatrix} dx + \begin{bmatrix} \nabla_{zx} f_1 & & \\ & \ddots & \\ & & \nabla_{zx} f_L \end{bmatrix} \begin{bmatrix} dz_0 \\ \vdots \\ dz_{L-1} \end{bmatrix}$$

Each vector dz_ℓ in turn varies with dx via $dz_\ell = (\nabla_x f_\ell) dx_\ell + (\nabla_z f_\ell) dz_{\ell-1}$, with the base case $dz_0 = 0$, since z_0 is the input and it does not vary with dx . Stacking up this recurrence gives

$$\underbrace{\begin{bmatrix} I & & \\ -\nabla_z f_2 & I & \\ & & \ddots \\ & & -\nabla_z f_{L-1} & I \end{bmatrix}}_{\equiv N} \begin{bmatrix} dz_0 \\ dz_1 \\ \vdots \\ dz_{L-1} \end{bmatrix} = \underbrace{\begin{bmatrix} 0 & & \\ \nabla_x f_1 & & \\ & \ddots & \\ & & \nabla_x f_{L-1} \end{bmatrix}}_{\equiv D^\vee} dx.$$

Plugging back the vector of dz 's gives

$$\begin{bmatrix} \text{vec}(dD_1) \\ \vdots \\ \text{vec}(dD_L) \end{bmatrix} = \left(\begin{bmatrix} \nabla_{xx} f_1 & & \\ & \ddots & \\ & & \nabla_{xx} f_L \end{bmatrix} + \begin{bmatrix} \nabla_{zx} f_1 & & \\ & \ddots & \\ & & \nabla_{zx} f_L \end{bmatrix} N^{-1} D^\vee \right) dx.$$

Plugging back this vector of dD 's gives

$$(dD)b = \begin{bmatrix} b_1^\top \otimes I & & \\ & \ddots & \\ & & b_L^\top \otimes I \end{bmatrix} \left(\begin{bmatrix} \nabla_{xx} f_1 & \cdots & \\ & \ddots & \\ & & \nabla_{xx} f_L \end{bmatrix} + \begin{bmatrix} \nabla_{zx} f_1 & \cdots & \\ & \ddots & \\ & & \nabla_{zx} f_L \end{bmatrix} N^{-1} D^\searrow \right) dx.$$

The term involving dM

The matrix dM has entries dM_2, \dots, dM_L on its off-diagonal blocks, and is zero everywhere else. Similar to the above, we can write

$$DM^{-1}(dM)b = DM^{-1} \begin{bmatrix} 0 & dM_2 & & \\ & 0 & dM_3 & \\ & & \ddots & \\ & & & 0 & dM_L \end{bmatrix} b \quad (13)$$

$$= DM^{-1} \begin{bmatrix} (dM_2)b_2 \\ \vdots \\ (dM_L)b_L \end{bmatrix} \quad (14)$$

$$= DM^{-1} \begin{bmatrix} (b_2^\top \otimes I) \text{vec}(dM_2) \\ \vdots \\ (b_L^\top \otimes I) \text{vec}(dM_L) \end{bmatrix} \quad (15)$$

$$= DM^{-1} \begin{bmatrix} 0 & b_2^\top \otimes I & & \\ & 0 & b_3^\top \otimes I & \\ & & \ddots & \\ & & & 0 & b_L^\top \otimes I \end{bmatrix} \begin{bmatrix} \text{vec}(dM_1) \\ \vdots \\ \text{vec}(dM_L) \end{bmatrix}. \quad (16)$$

Each entry $dM_\ell = d\nabla_z f_\ell(z_{\ell-1}; x_\ell)$ varies with dx through both x_ℓ and $z_{\ell-1}$ as $d\text{vec}M_\ell = \nabla_{xz} f_\ell dx_\ell + \nabla_{zz} f_\ell dz_{\ell-1}$. Following the steps of the previous section, we can write

$$\begin{bmatrix} \text{vec}(dD_1) \\ \vdots \\ \text{vec}(dD_L) \end{bmatrix} = \left(\begin{bmatrix} \nabla_{xx} f_1 & \cdots & \\ & \ddots & \\ & & \nabla_{xx} f_L \end{bmatrix} + \begin{bmatrix} \nabla_{zx} f_1 & \cdots & \\ & \ddots & \\ & & \nabla_{zx} f_L \end{bmatrix} N^{-1} D^\searrow \right) dx.$$

Plugging back gives

$$DM^{-1}(dM)b = DM^{-1} \begin{bmatrix} 0 & b_2^\top \otimes I & & \\ & 0 & b_3^\top \otimes I & \\ & & \ddots & \\ & & & 0 & b_L^\top \otimes I \end{bmatrix} \left(\begin{bmatrix} \nabla_{xx} f_1 & \cdots & \\ & \ddots & \\ & & \nabla_{xx} f_L \end{bmatrix} + \begin{bmatrix} \nabla_{zx} f_1 & \cdots & \\ & \ddots & \\ & & \nabla_{zx} f_L \end{bmatrix} N^{-1} D^\searrow \right) dx.$$

Putting it all together

We have just shown that the Hessian of the deep net has the form

$$H \in \mathbb{R}^{pL \times pL} \equiv \frac{\partial^2 z_L}{\partial x^2} = D_b (D_{xx} + D_{zx} N^{-1} D_x^\searrow) + D_x M^{-1} B_b (D_{xz} + D_{zz} N^{-1} D_x^\searrow).$$

The definitions below annotates the size of the various matrices in this expression assuming the first layer has a -dimensional activations ($z_1 \in \mathbb{R}^a$) and p -dimensional parameters ($x_\ell \in \mathbb{R}^p$). The backward variable $b_1 \in \mathbb{R}^a \equiv \frac{\partial z_L}{\partial z_\ell}^\top$ is computed during the backward pass:

$$\begin{aligned}
D_b &\equiv \begin{bmatrix} \underbrace{b_1^\top \otimes I}_{p \times ap} & & \\ & \ddots & \\ & & b_L^\top \otimes I \end{bmatrix}, B_b \equiv \begin{bmatrix} 0 & \underbrace{b_2^\top \otimes I}_{a \times a^2} & & \\ & 0 & b_3^\top \otimes I & \\ & & \ddots & \\ & & & 0 & b_L^\top \otimes I \end{bmatrix}, \\
D_x &\equiv \begin{bmatrix} \underbrace{\nabla_x f_1}_{a \times p} & & \\ & \ddots & \\ & & \nabla_x f_L \end{bmatrix}, D_x^\searrow \equiv \begin{bmatrix} 0 & & & \\ \nabla_x f_1 & & & \\ & \ddots & & \\ & & \nabla_x f_{L-1} & \end{bmatrix}, \\
D_{xx} &\equiv \begin{bmatrix} \underbrace{\nabla_{xx} f_1}_{ap \times p} & & \\ & \ddots & \\ & & \nabla_{xx} f_L \end{bmatrix}, D_{xz} \equiv \begin{bmatrix} \underbrace{\nabla_{xz} f_1}_{a^2 \times p} & & \\ & \ddots & \\ & & \nabla_{xz} f_L \end{bmatrix}, \\
D_{zx} &\equiv \begin{bmatrix} \underbrace{\nabla_{zx} f_1}_{ap \times a} & & \\ & \ddots & \\ & & \nabla_{zx} f_L \end{bmatrix}, D_{zz} \equiv \begin{bmatrix} \underbrace{\nabla_{zz} f_1}_{a^2 \times a} & & \\ & \ddots & \\ & & \nabla_{zz} f_L \end{bmatrix}, \\
M &\equiv \begin{bmatrix} I & \underbrace{-\nabla_z f_2}_{a \times a} & & \\ & I & -\nabla_z f_3 & \\ & & \ddots & \\ & & & I & -\nabla_z f_L \\ & & & & 1 \end{bmatrix}, N \equiv \begin{bmatrix} I & & & \\ \underbrace{-\nabla_z f_2}_{a \times a} & I & & \\ & & \ddots & \\ & & & -\nabla_z f_{L-1} & I \end{bmatrix}.
\end{aligned}$$

While all of the matrices involved in the Hessian are either diagonal or bi-diagonal, the Hessian itself is dense because some of these matrices are inverted (in particular, M and N).

The inverse of the Hessian

$$\begin{aligned}
H &= D_b (D_{xx} + D_{zx} N^{-1} D_x^\searrow) + D_x M^{-1} B_b (D_{xz} + D_{zz} N^{-1} D_x^\searrow) \\
&= D_b D_{xx} + D_b D_{zx} N^{-1} D_x^\searrow + D_x M^{-1} B_b D_{xz} + D_x M^{-1} B_b D_{zz} N^{-1} D_x^\searrow
\end{aligned}$$

Convergence Rate

Above, I claimed that a good preconditioner is one that causes the gradient to point toward the minimizer. This section derives a bound on the convergence rate of stochastic gradient descent in terms of this alignment.